

Deep dive into Robotics I

AI: Algorithms to Actuation

Who is Kv?

Why should you hear me?


Python Programming

ROBOPACK

Powerful **multi-purpose** programming language

Simple and **easy-to-use** syntax


Most popular first-choice language for beginners!



By definition!

Python is an **interpreted, object-oriented, high-level** programming language. As it is general-purpose, it has a wide range of applications from web development, building desktop GUI to scientific and mathematical computing.

Applications : Web Applications



Hello, World! – a breakdown!

ROBOPACK

```
print("Hello, World!")
```

Here, `print()` is a function tells the computer to perform an action. What is that action? To output the value residing in it!

Hello, World! – a breakdown!



```
print("Hello, World!")
```

Inside the parentheses of the `print()` function is a sequence of characters — `Hello, World!` — that is enclosed in quotation marks `'` or `"`.

Any characters that are inside of quotation marks are called a *string*.

Data types

Every value in Python has a **datatype**. Since everything is an object in Python programming, data types are actually **classes** and **variables** are instance (object) of these classes.

List

Ordered sequence of items.

```
1 x = [6, 99, 77, 'Apple']  
2 print(x, "is of type", type(x))
```

```
[6, 99, 77, 'Apple'] is of type <class 'list'>
```


Dictionary

Unordered collection of **key-value** pairs. Optimized for **retrieving data**.

```
1 d = {1: 'Apple', 2: 'Cat', 3: 'Food'}  
2 print(d, type(d))  
3  
4 d[3]
```

```
{1: 'Apple', 2: 'Cat', 3: 'Food'} <class 'dict'>  
'Food'
```

Variables



To *temporarily* store data in the computer's memory, for example, someone's name, age, email, address, price of a product, and so on.

Arithmetic Operators

Perform **mathematical operations** like addition, subtraction, multiplication etc.

Symbol	Task Performed	Meaning	Example
+	Addition	add two operands or unary plus	$x + y$ or $+2$
-	Subtraction	subtract right operand from the left or unary minus	$x - y$ or -2
*	Multiplication	Multiply two operands	$x * y$
/	Division	Divide left operand by the right one (always results into float)	x / y
%	Modulus (remainder)	remainder of the division of left operand by the right	$x \% y$ (remainder of x/y)
//	Integer/Floor division	division that results into whole number adjusted to the left in the number line	$x // y$
**	Exponentiation (power)	left operand raised to the power of right	$x ** y$ (x to the power y)

Comparison Operators

Used to **compare values**. It either returns **True** or **False** according to the condition.

Symbol	Task Performed	Meaning	Example
>	greater than	True if left operand is greater than the right	$x > y$
<	less than	True if left operand is less than the right	$x < y$
==	equal to	True if both operands are equal	$x == y$
!=	not equal to	True if both operands are not equal	$x != y$
>=	greater than or equal to	True if left operand is greater than or equal to the right	$x >= y$
<=	less than or equal to	True if left operand is less than or equal to the right	$x <= y$

Logical Operators

Symbol	Meaning	Example
and	True if both the operands are true	x and y
or	True if either of the operand is true	x or y
not	True if operand are false (complements the operand)	not x

Functions



block of **organized**, **reusable** (DRY- Don't Repeat Yourself) code with a name that is used to perform a single, specific task.

```
def function_name(parameter1, parameter2):  
    """docstring"""  
    # function body  
    # write some action  
    return value
```




```
1 # Example 1:  
2 def greet():  
3     print("Welcome to Python for Data Science")  
4  
5 # call function using its name  
6 greet()
```



Welcome to Python for Data Science

Difference between AI/ML/DL

ROBOPACK



Machine Learning

Traditional Programming



Machine Learning



branch of artificial intelligence (AI) and computer science that focuses on the using **data** and **algorithms** to enable AI to *imitate the way that humans learn*, gradually improving its accuracy.

Supervised Learning


Given training data

+

Desired output

=

Predict

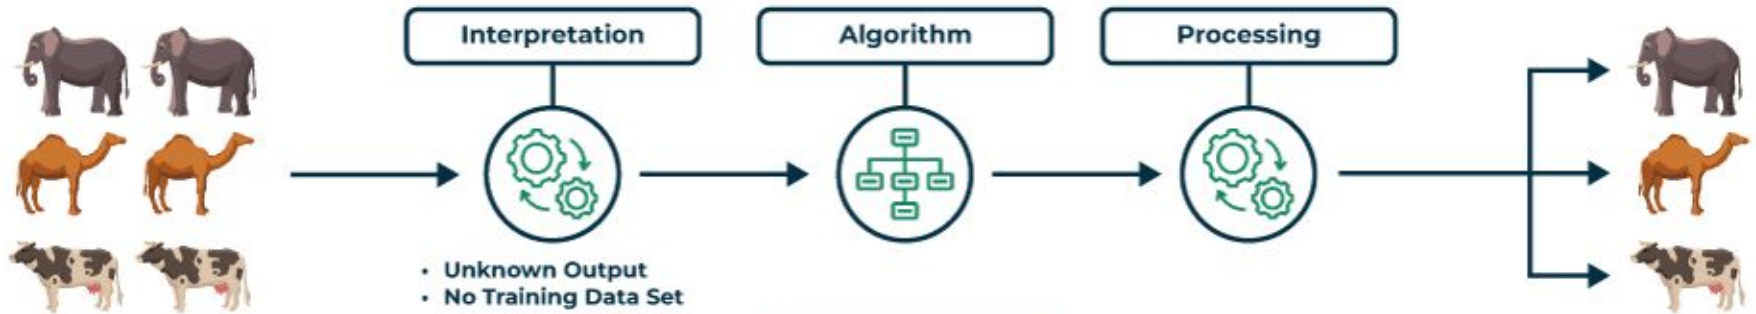


Unsupervised Learning

Lots of Training data (w/o desired output)

=

Find patterns



Reinforcement Learning

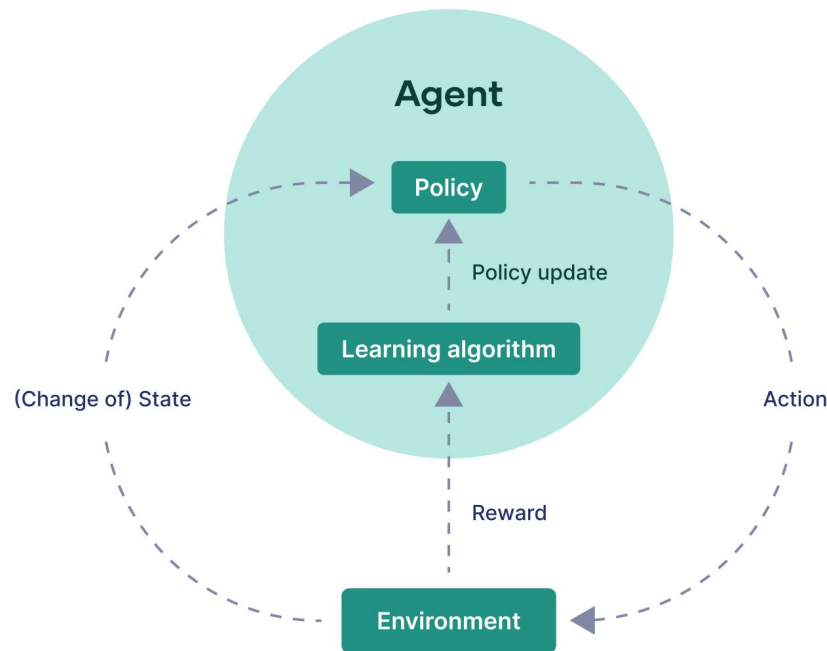
trains software to make decisions to achieve the most optimal results

almost no data
(just game parameters)

=

excel in the game

ROBOPACK




Let's compare!

Supervised Learning

Data: (x, y)
x is data, y is label

Goal:
Learn function to map
 $x \rightarrow y$

Example:




It's a ['mango']

Unsupervised Learning

Data: x
x is data, no labels!

Goal:
Learn underlying structure

Example:




This thing is like the other

Reinforcement Learning

Data:
state-action pair

Goal:
maximize future rewards over
many time stamps

Example:



Eat it! 'coz that's wot u live for




Figure 3.1: The agent–environment interaction in reinforcement learning.

Reinforcement Learning




Agent: takes action.

Reinforcement Learning



Environment: world in which the agent exists (or operates).

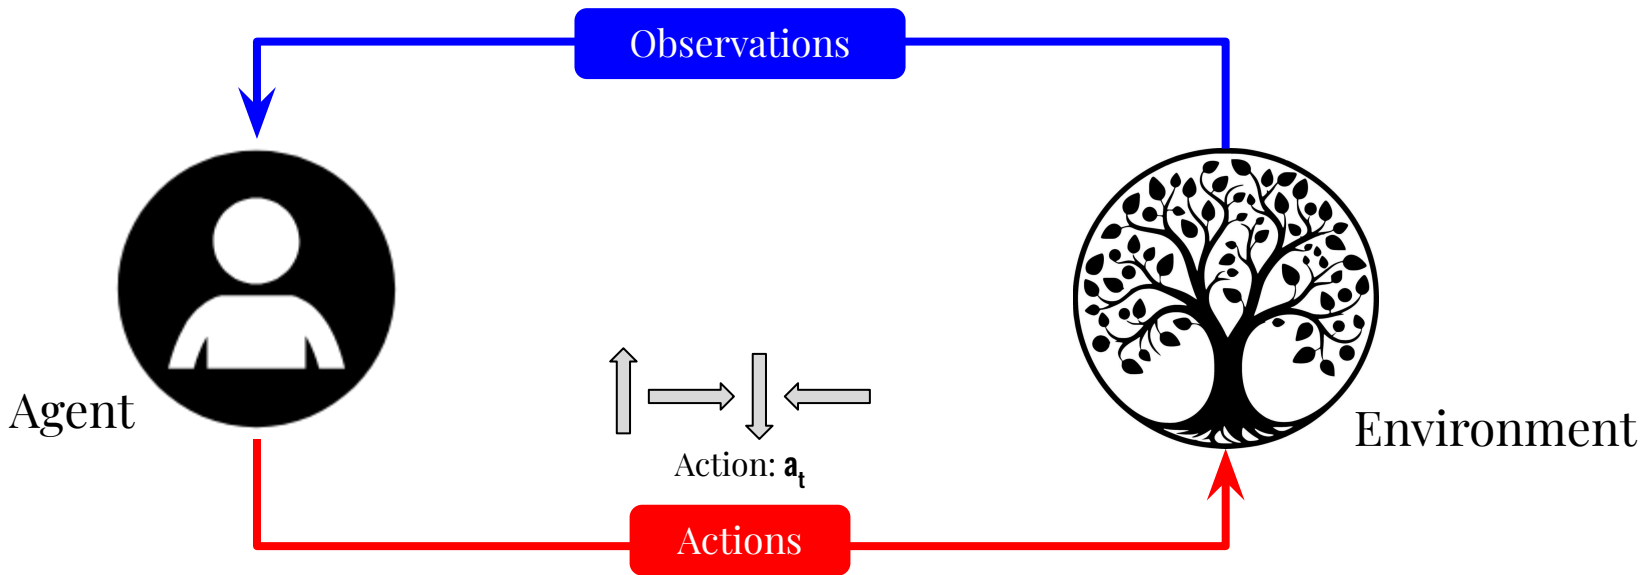
Reinforcement Learning



Action: a move that agent can make in the environment.

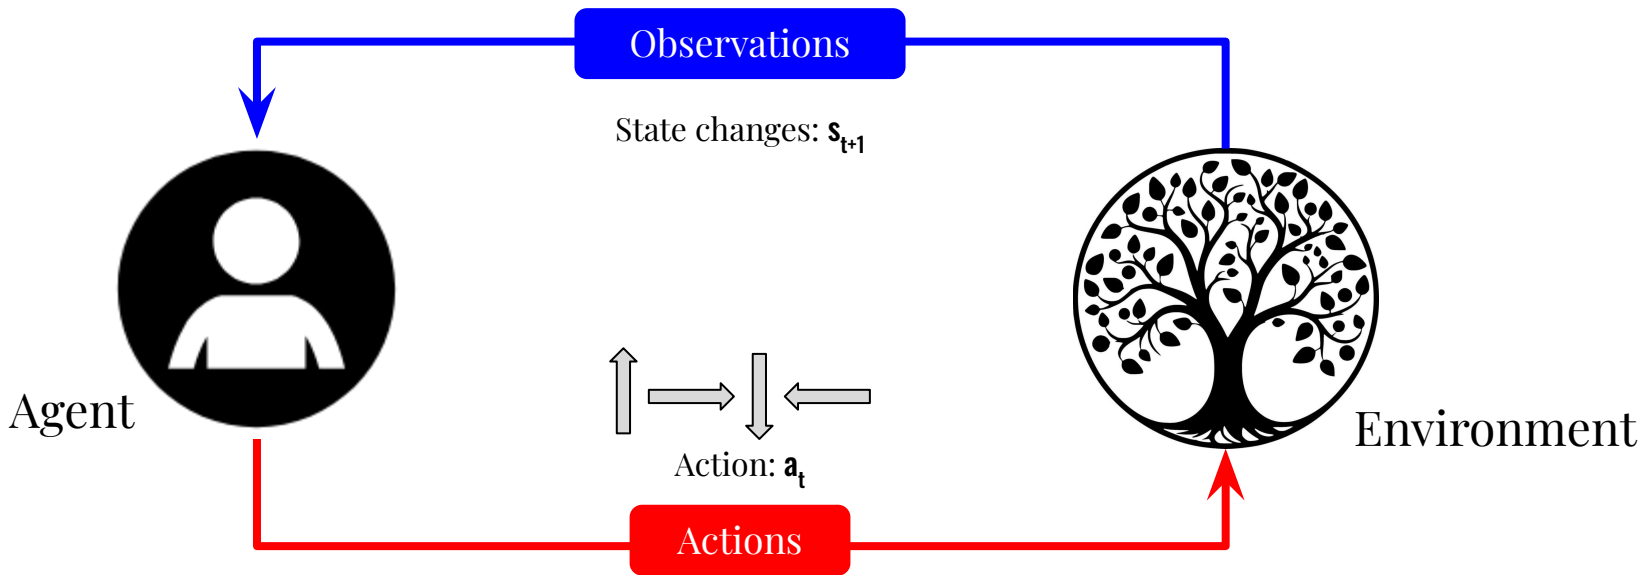
Action space A : a set of possible actions an agent can make in the environment.

Reinforcement Learning



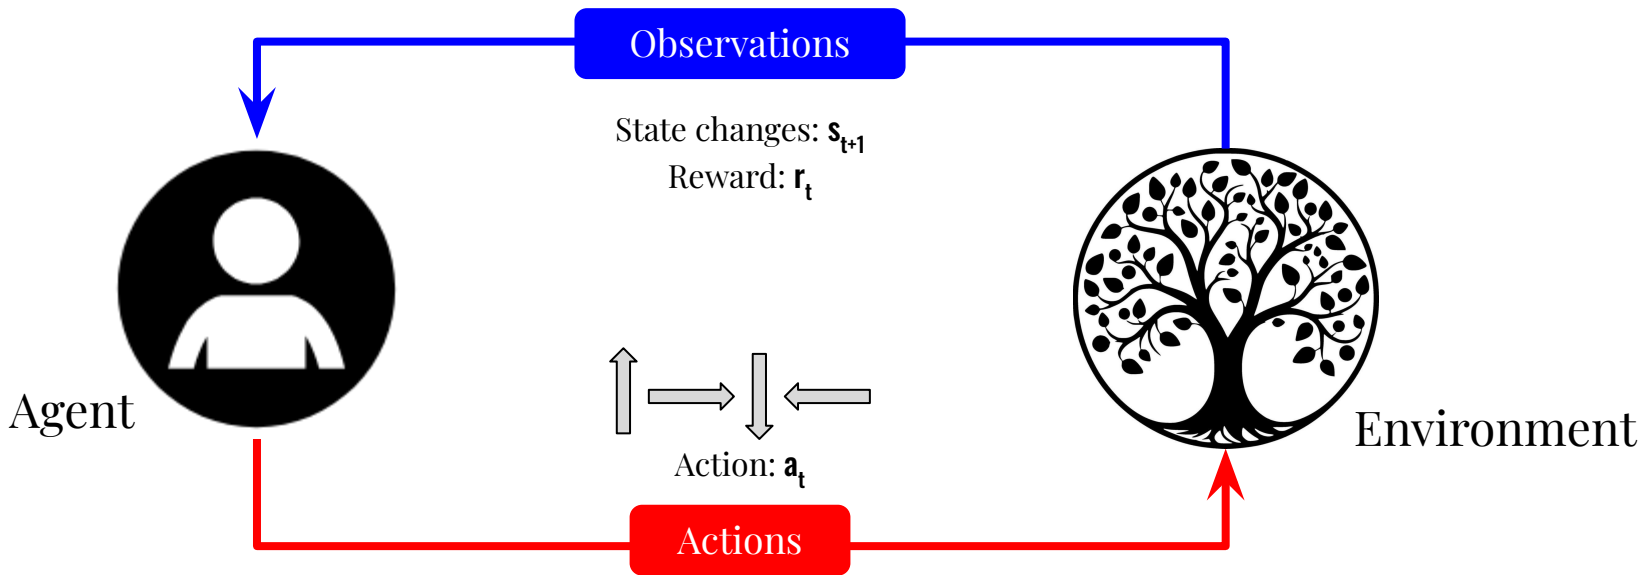
Observations: of the environment after taking actions.

Reinforcement Learning




State: a situation that the agent perceives.

Reinforcement Learning



Reward: feedback that measures the success/failure of the agent's action.


Reinforcement Learning



Total Reward
(Return)

$$R_t = \sum_{i=t}^{\infty} r_i$$


Reinforcement Learning



Total Reward
(Return)

$$R_t = \sum_{i=t}^{\infty} r_i = r_t + r_{t+1} + \dots + r_{t+n} + \dots$$


Reinforcement Learning



Discounted
Total Reward
(Return)

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i$$

Reinforcement Learning



Discounted
Total Reward
(Return)

$$R_t = \sum_{i=t}^{\infty} \gamma^i r_i = \gamma^t r_t + \gamma^{t+1} r_{t+1} + \dots + \gamma^{t+n} r_{t+n} + \dots$$

Defining the Q-function

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Total reward, R_t is the discounted sum of all rewards obtained from time t .

$$Q(s_t, a_t) = E[R_t \mid s_t, a_t]$$

The Q-function captures the expected total future reward an agent in state, \mathbf{s} , can perceive by executing a certain action, \mathbf{a}

How to take action given a Q-function?

$$Q(s_t, a_t) = E[R_t \mid s_t, a_t]$$

(state, action)

The agent needs a **policy** $\pi(s)$, to infer the **best action to take** at its state, s

Strategy: the policy should choose an action that maximizes future reward

$$\pi^*(s) = \arg \max_a Q(s, a)$$

Deep Reinforcement Learning Algorithms

Value Learning

Find $Q(s,a)$

$$a = \underset{a}{\operatorname{argmax}} Q(s,a)$$

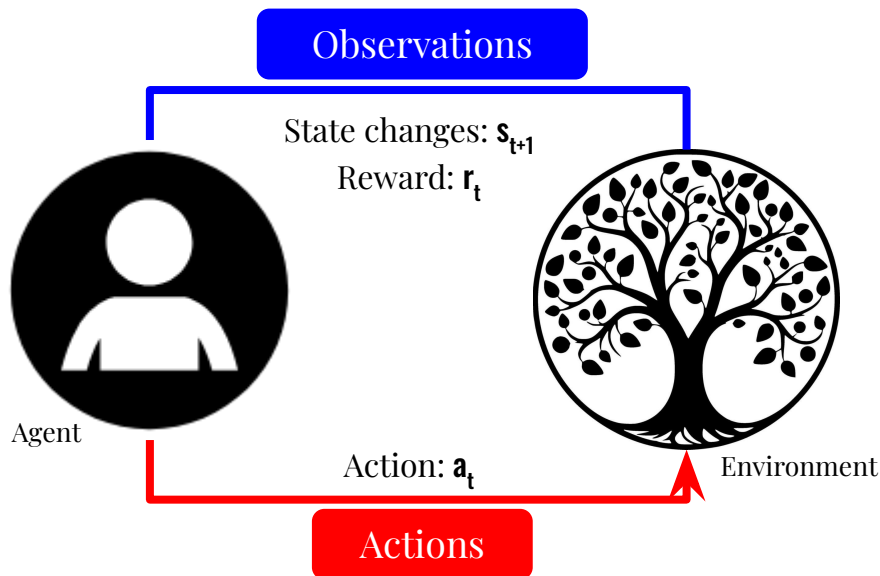
Policy Learning

Find $\pi(s)$

Sample $a \sim \pi(s)$

Training Policy Gradient

Reinforcement Learning Loop



Case Study – Self Driving Car

Agent: vehicle

State: camera, lidar, etc

Action: steering wheel angle

Reward: distance traveled

Get in touch!



<https://www.linkedin.com/in/keivalya/>



@keivalya



keivalya.pandya@bvmengineering.ac.in